
Hierarchical Lateral Inhibition Enables Robust Winner-Take-All Computation in Large-Scale Cortical Models

Anonymous Author(s)

Affiliation

Address

email

Abstract

In computational neuroscience, models of visual cortex have long captured classical receptive field properties, yet the circuit mechanisms underlying extra-classical surround suppression remain unresolved. In vivo experiments have consistently shown a winner-take-all (WTA) dynamic, where a strongly stimulated cortical column suppresses neighboring activity by over 75%; this dynamic is essential for basic visual functions like figure-ground segmentation. Classical models attribute this computation to lateral inhibition within V1, yet prior work has struggled to reproduce the magnitude of suppression observed physiologically using models of V1 circuits alone. In this work, we hypothesized that robust WTA emergence requires extending lateral inhibitory competition to higher cortical areas, beyond V1. To test this hypothesis, we simulated a large-scale spiking neural network based on the connectivity of the canonical cortical microcircuit and systematically compared three different architectures: (1) V1 columns with lateral inhibition restricted to V1, (2) a hierarchical V1-V2 model with simple feedback between V1 and V2 but no lateral inhibition in V2, and (3) a hierarchical model with feedback where both V1 and V2 incorporate lateral inhibitory connections. Our experiments show that only when lateral competition is extended to V2 do the circuits generate a physiologically accurate surround suppression effect, achieving a suppression index of 75% that quantitatively matches in vivo recordings. In this architecture, V2 lateral inhibition gates the excitatory feedback returning to V1, enabling the network to dynamically encode relative stimulus strength even when absolute input remains constant. Our work explains a long-standing physiological puzzle and establishes that lateral inhibition in higher cortical areas is necessary to drive the circuit from a regime of general divisive normalization into a robust winner-take-all (WTA) state, which is essential for strong surround suppression. Moreover, this hierarchical inhibitory architecture provides a blueprint for designing systems that require context-dependent information processing.

1 Introduction

From a computational point of view, the contemporary understanding of sensory processing in the visual system is that the cortex is organized as a hierarchy of filters where neurons extract progressively complex features of the visual scene through stage-by-stage processing. Central to this conception is the classical receptive field (CRF), operationally defined as the specific region of visual space where the presence of a stimulus, such as light or a dark bar, can evoke a direct neuronal response, increasing the neuron's firing rate. In this classical model, processing has been interpreted as a local phenomenon, where the response of each functional unit is determined almost exclusively by the physical properties of the stimulus and its location within the field of view.

37 However, the strictly local view of neuronal processing has been challenged by the discovery of
 38 extra-classical receptive field (eCRF) effects. In this phenomenon, sensory responses are not only
 39 determined by stimuli within the CRF but are also robustly modulated by contextual information
 40 from the surrounding visual field [1]. The most prominent of these effects is surround suppression,
 41 a non-linear phenomenon in which expanding a stimulus beyond the CRF leads to a substantial
 42 decrease in firing rates. This implies a high level of global integration, challenging the conception
 43 of primary visual cortex as merely a set of isolated local filters, and suggests that the functional
 44 architecture of the visual system must incorporate mechanisms capable of integrating information
 45 across disparate spatial locations [2]. Long-range horizontal connections within V1 were initially
 46 proposed as the primary substrate for eCRF effects; however, accumulating evidence indicates that
 47 these lateral pathways are physiologically limited in both conduction velocity and spatial reach [3, 4].
 48 This has led to the prevailing hypothesis that feedback projections from higher-order areas, such as
 49 V2, play a decisive role in integrating global information, providing the necessary spatial extent and
 50 conduction speed to modulate V1 activity effectively [5].

51 Despite significant advances in developing biologically plausible cortical microcircuits, such as the
 52 canonical model proposed by Potjans and Diesmann [6], a substantial gap remains between theoretical
 53 predictions and experimental observations of eCRF effects. Most existing computational models fail
 54 to quantitatively replicate the magnitude of surround suppression observed physiologically, which
 55 can reach a reduction of approximately 75% in neuronal firing rates [7, 8]. This discrepancy often
 56 stems from feedback implementations that are either linear or excessively simplified. By neglecting
 57 the laminar architecture of the cortex and the specificity of inter-areal connections, these models fail
 58 to capture the intrinsic competitive dynamics of the visual hierarchy.

59 In this study, we hypothesized that robust surround suppression is not merely a byproduct of hi-
 60 erarchical feedback but an emergent property of competitive dynamics within the visual system.
 61 Utilizing the NEST simulator [9], we implemented a large-scale spiking neural network model
 62 based on the canonical cortical microcircuit to examine how inter-areal interactions shape eCRF
 63 effects. Specifically, we propose that lateral competition within higher-order areas, particularly V2, is
 64 the essential mechanism for regulating top-down modulation and achieving the suppression levels
 65 observed experimentally. By comparing multiple architectural configurations, we aim to demonstrate
 66 that inter-areal competition is the missing component required to replicate the non-linear contextual
 67 dynamics of the primary visual cortex.

68 Our main contributions are as follows. First, we systematically compare three hierarchical architec-
 69 tures of large-scale spiking cortical microcircuits, isolating the specific role of lateral inhibition in V2
 70 for extra-classical surround suppression. Second, we demonstrate that lateral competition within V2
 71 is necessary to achieve the 75% suppression index consistently reported in primate in vivo record-
 72 ings; architectures relying solely on V1 lateral connections or non-competitive V2 feedback both
 73 fail to replicate this magnitude. Finally, we show that this hierarchical winner-take-all mechanism
 74 enables the network to encode relative stimulus strength even when absolute input remains constant,
 75 revealing a computational principle directly transferable to artificial perception systems that require
 76 context-dependent normalization.

77 2 Related Work

78 Early computational approaches to extra-classical receptive field (eCRF) effects primarily emphasized
 79 the role of intrinsic horizontal connections within the primary visual cortex (V1) [10, 11]. These
 80 lateral projections enable spatial integration by linking neurons with non-overlapping receptive fields
 81 across the cortical surface. However, horizontal connections face significant physiological constraints
 82 in both spatial extent and conduction velocity, which limit their capacity to account for the rapid,
 83 long-range dynamics of surround suppression observed experimentally [3, 4]. Consistent with these
 84 biological constraints, simulations of lateral inhibition restricted to V1 yield only modest suppression,
 85 typically a 30–35% reduction in firing rate [12, 13], which is far below the 75% reduction reported
 86 in vivo. These findings suggest that local V1 circuitry alone lacks the dynamic range necessary to
 87 fully account for global contextual integration.

88 An alternative hypothesis proposes that feedback projections from higher-order areas, particularly V2,
 89 serve as the primary driver of long-range contextual modulation [5]. However, existing computational
 90 implementations of this feedback have largely relied on simplified loops that treat top-down signals

as linear additive inputs or non-specific gain modulators [14, 15]. Such formulations fail to capture the highly non-linear and stimulus-specific nature of eCRF effects. Critically, no existing model has incorporated lateral inhibitory competition within the feedback-generating higher-order areas—a mechanism that may be essential for achieving the precise spatial and feature-specific filtering observed physiologically. This gap motivates the present study: we investigate whether embedding competitive dynamics within V2 feedback pathways can yield the robust surround suppression that purely local or linearly integrated models have failed to reproduce.

Substantial progress has been made in developing biologically constrained spiking neural network models of cortical processing. The canonical microcircuit model of Potjans and Diesmann [6] has become a widely adopted framework for reproducing the laminar connectivity and asynchronous-irregular dynamics of local cortical patches. This model has been successfully extended to multi-area hierarchies [16] and applied to cognitive tasks such as working memory and decision-making [17]. Concurrently, spiking networks trained via surrogate gradients have achieved competitive performance on visual recognition benchmarks [18], demonstrating that spike-based computation can scale to complex tasks. However, these models typically treat hierarchical interactions as feedforward feature extraction or simple convergent feedback, without incorporating lateral inhibitory competition within higher areas. Whether the addition of such competitive dynamics would enable these networks to exhibit the robust, stimulus-specific surround suppression characteristic of biological vision remains an open question, which we address in this work.

The phenomenon of surround suppression has been extensively studied through the lens of divisive normalization, a canonical neural computation in which a neuron’s response is divided by the aggregate activity of a broader neuronal pool [15]. Early formulations by Heeger [19] demonstrated that divisive normalization can account for contrast-dependent saturation and cross-orientation suppression in V1 using simple computational rules. Subsequent work generalized this principle to a wide range of sensory and cognitive phenomena [20, 21]. In parallel, the machine learning community has independently converged on normalization as an essential architectural component. Batch normalization [22], layer normalization [23], and divisive normalization networks [24] have each shown that rescaling activations relative to local or global statistics improves training stability and generalization. Importantly, these implementations operate through feedforward or self-normalizing mechanisms within a single processing stage, and do not incorporate the inter-areal competitive dynamics that characterize cortical hierarchies.

In this work, we propose a hierarchical competition mechanism to explain eCRF effects. This mechanism is closely aligned with predictive coding frameworks [25, 26], in which higher cortical areas generate top-down predictions that are compared against incoming sensory evidence from lower levels; mismatches produce prediction errors that propagate upward, while accurately predicted features are suppressed through a process known as *explaining away*. This computation bears functional similarities to biased competition models of attention [27, 28], where top-down signals resolve competition among multiple stimuli by selectively enhancing or suppressing specific feature representations. The winner-take-all competition that we implement in V2 can be interpreted as a circuit-level instantiation of these principles: when a peripheral stimulus drives a competing V2 module, it actively inhibits the center-aligned module’s feedback, effectively signaling that the central activity is redundant given the broader spatial context. By grounding these abstract theories in a concrete spiking microcircuit with laminar specificity, our work bridges normative computational principles and their biologically plausible implementation.

In summary, none of these research lines, including biophysical models of cortical microcircuits, divisive normalization in theory and practice, and predictive coding, explain how extra-classical surround suppression emerges from hierarchical interactions. Specifically, the role of lateral inhibitory competition in higher visual areas as a gating mechanism for top-down feedback remains unexplored in large-scale spiking models. Our work addresses this gap by demonstrating that V2 lateral inhibition is necessary, and within the architectures tested sufficient, to reproduce the hallmark signatures of robust sensory normalization.

3 Methods

We implemented a large-scale spiking neural network using the NEST simulator [9]. We constructed three architectural configurations of increasing complexity, all grounded in the canonical cortical

microcircuit [6]: (1) V1 columns with only lateral intra-areal connections, (2) a V1-V2 hierarchy with simple convergent feedback, and (3) a V1-V2 hierarchy with reciprocal lateral inhibition between V2 modules. In the following sections we describe the neuron model, the laminar microcircuit architecture, the hierarchical connectivity, the stimulation protocols, and the data analysis pipeline.

3.1 Neuron Model and Simulation Environment. Neurons were modeled as leaky integrate-and-fire (LIF) units with exponential-decaying postsynaptic currents. The subthreshold membrane potential dynamics are governed by the membrane time constant ($\tau_m = 10$ ms), membrane capacity ($C_m = 250$ pF), and resting potential ($V_{\text{rest}} = -65$ mV). A spike is emitted when V_m reaches the threshold ($V_{\text{th}} = -50$ mV), after which the potential is clamped to the reset value ($V_{\text{reset}} = -65$ mV) for an absolute refractory period ($\tau_{\text{ref}} = 2$ ms). Excitatory synaptic weights were set to $w = 87.8 \pm 8.8$ pA with a decay time constant $\tau_{\text{syn}} = 0.5$ ms, and relative inhibitory synaptic strength was fixed at $g = -4$. Transmission delays were drawn from uniform distributions of 1.5 ± 0.75 ms for excitatory synapses and 0.75 ± 0.38 ms for inhibitory synapses. Background cortical activity was emulated via Poisson spike trains delivered to all neurons at a mean rate of 8 Hz. All parameters were tuned to sustain the asynchronous-irregular firing regime characteristic of awake cortical states. Simulations were performed using the NEST simulator [9]. The membrane potential evolves according to

$$\tau_m \frac{dV_m}{dt} = -(V_m - V_{\text{rest}}) + R_m I_{\text{syn}}(t), \quad (1)$$

where R_m is the membrane resistance and $I_{\text{syn}}(t)$ is the total synaptic current. Each synaptic input generates a conductance-based current with exponential decay,

$$I_{\text{syn}}(t) = w e^{-t/\tau_{\text{syn}}}, \quad (2)$$

where w is the synaptic weight and τ_{syn} is the synaptic time constant.

3.2 Laminar Microcircuit Architecture. Each V1 module represents a 1 mm^2 patch of cortical tissue, based on the full-scale canonical microcircuit model proposed by Potjans and Diesmann [6]. It comprises approximately 77,000 neurons distributed across four laminar layers: L2/3, L4, L5, and L6. Within each layer, the network maintains a physiological ratio of excitatory to inhibitory neurons of approximately 80/20. Local recurrent connectivity follows empirically derived connection probabilities.

To capture intra-areal spatial integration, long-range horizontal projections within V1 are incorporated. These lateral connections are restricted to the superficial L2/3 layer and are functionally segregated by distance. Nearby connections link functionally similar adjacent columns with both excitatory and inhibitory synapses, with transmission delays of $d_{\text{n,lat}} = 5 \pm 1.5$ ms. Distant connections, in contrast, link columns with disparate orientation preferences and operate via a predominantly suppressive mechanism through disinaptic inhibition, with a slower conduction delay of $d_{\text{d,lat}} = 8 \pm 2.5$ ms (see **Figure 2 (Top)**).

The complete connection probability matrix between the eight neuronal populations (four layers $\times$ two types) follows the values reported by Potjans and Diesmann [6] and is provided in Table 3 of the Appendix.




Figure 1: Canonical Microcircuit and Hierarchical Architecture. (Left) Local and intra-areal connectivity of the V1 microcircuit, detailing the specific interlaminar excitatory (E) and inhibitory (I) projections. (Right) Inter-areal feedforward and feedback connectivity between V1 and V2, along with near and distant lateral connections within V1.

3.3 Hierarchical Integration and Competitive Dynamics. To investigate the mechanistic origins of eCRF effects, the baseline architecture was extended to include a higher-order visual area (V2). The V2 module mirrors the cytoarchitectural variations in laminar thickness and neuronal densities reported for the macaque V2 region. Inter-areal communication is mediated by topographic projections with fast transmission delays of 2 ± 0.5 ms. Feedforward pathways originate from the superficial layers L2/3 of V1 and terminate predominantly in L4 and L6 of V2. Feedback projections originate from the deep layers L5 and L6 of V2 and target the superficial L2/3 and infragranular L5 of V1, anatomically avoiding L4 to provide modulatory rather than driving input (see Figure 1 (Right)).

To test the hypothesis that eCRF phenomena emerge from inter-areal interactions rather than purely local filtering, we embedded a lateral competition mechanism within the V2 module. This architecture comprises two spatially distinct, feature-aligned V2 modules that engage in competition via strong, reciprocal inhibitory connections. This hierarchical winner-take-all dynamic non-linearly refines the top-down feedback signals before they modulate V1 activity, overcoming the limitations of linear feedback models and enabling the robust, stimulus-specific surround suppression observed in physiological recordings.

To isolate the contribution of each circuit component, we systematically compared three architectural configurations (see Figure 1). **(1) V1 lateral only:** Four V1 microcolumns arranged in a ring, interconnected exclusively via horizontal connections within L2/3. No V2 modules are present. **(2) V1 + linear V2 feedback:** The V1 columns project feedforward to a single V2 module, which pools inputs from all columns and returns non-specific feedback to V1. This configuration implements hierarchical feedback without lateral competition. **(3) V1 + competitive V2 feedback:** The V1 columns project feedforward to two V2 modules, each aligned to a specific orientation column. The two V2 modules are interconnected via strong, reciprocal inhibitory synapses, implementing a winner-take-all competition. All other parameters, including neuron counts, local connectivity, and stimulation protocols, were held constant across configurations to ensure that differences in suppression magnitude can be attributed exclusively to the architectural manipulation.

3.4 Sensory Stimulation Protocols. Bottom-up sensory stimulation was emulated via thalamocortical afferents originating from the lateral geniculate nucleus (LGN), projecting primarily to layers L4 and L6 of V1. Thalamic inputs were modeled as Poisson spike trains. To simulate the classical receptive field response, the thalamic firing rate directed at the central microcolumn was increased to 20 Hz after a 500 ms baseline period. To evaluate extra-classical modulation, a peripheral contextual stimulus was applied to a distant column at 1000 ms. The intensity of this peripheral stimulus was systematically varied across three levels: 10, 20, and 30 Hz (Figure 2).

To quantify the strength of surround suppression, we define the suppression index (SI) as

$$SI = 1 - \frac{R_{\text{center+surround}}}{R_{\text{center}}}, \quad (3)$$

where R_{center} is the mean firing rate of the central column during classical receptive field stimulation alone, and $R_{\text{center+surround}}$ is the mean firing rate during concurrent center and peripheral stimulation. An SI of 1 indicates complete suppression, and an SI of 0 indicates no contextual modulation.

3.5 Data Analysis. Network dynamics and mean firing rates were evaluated using peri-stimulus time histograms (PSTH). To connect the spiking network model to macroscopic electrophysiological biomarkers, including gamma-band oscillations and excitation/inhibition balance alterations, we simulated the local field potential (LFP) using a spatiotemporal kernel-based approximation. In this framework, the LFP was reconstructed as the linear superposition of unitary field potentials generated by individual action potentials, enabling spectral analysis of the network rhythmic activity under distinct contextual modulations.

4 Results

We present our findings in four parts. First, we validate the baseline dynamics and classical receptive field properties of the model. Second, we compare surround suppression across the three architectural configurations. Third, we demonstrate that only the competitive V2 architecture reproduces the $\sim 75\%$ suppression observed in vivo, and we confirm via virtual ablation that this effect is driven by hierarchical competition. Finally, we show that perturbing the excitation/inhibition balance replicates schizophrenia-linked biomarkers, establishing the translational relevance of the model.




Figure 2: **Stimulation Protocols and Network Configurations.** Schematic representation of the classical receptive field (CRF) and extra-classical receptive field (eCRF) paradigms across the three evaluated architectures. The bottom panels illustrate the thalamic stimulation over time and the corresponding peri-stimulus time histograms (PSTH) for Layer 2/3 neurons in V1, highlighting the varying degrees of surround suppression.

4.1 Spatial Summation and Classical Receptive Field Characterization. We first validated the stability of the full model in resting state and then characterized its classical receptive field responses. In the absence of targeted stimuli, the network maintained an asynchronous-irregular firing regime with stable baseline rates, averaging approximately 0.5 Hz in superficial layer L2/3 and 6.2 Hz in L5, confirming that the introduction of inter-columnar and inter-areal connectivity did not induce pathological oscillations. To characterize the classical receptive field, we evaluated the network response to variations in stimulus contrast, simulated by parametrically increasing the firing rate of the topographically aligned thalamic afferents from 5 to 30 Hz. The model exhibited a monotonic, contrast-dependent scaling of firing rates, most prominently in the primary thalamorecipient layer L4 and the superficial layer L2/3. To confirm spatial specificity, an optimal 20 Hz stimulus was delivered exclusively to the central microcolumn. While the stimulated central column responded vigorously, the spatially distant peripheral column remained at baseline firing levels. This demonstrates that local horizontal connections mediate effective lateral inhibition, preventing spurious propagation of strong sensory activation and preserving the spatial acuity of the classical receptive field.

4.2 Comparison of Architectural Configurations for Surround Suppression. To systematically evaluate the structural substrates of extra-classical receptive field effects, we tested the network under a dual-stimulus contextual protocol, in which a central stimulus was followed by an asynchronous peripheral stimulus of varying intensity. All reported values are means $\pm$ standard deviation across $n = 10$ simulation trials with different random seeds.

We compared the magnitude of surround suppression across the three architectural configurations. In the first configuration, with V1 lateral connections only, the peripheral stimulus induced only mild suppression. Under the highest contextual contrast (30 Hz), the L2/3 excitatory populations retained $63.3\% \pm 4.4\%$ of their peak classical receptive field activity, equivalent to approximately 36.7% suppression. L5 neurons showed virtually no suppressive modulation. This indicates that while intra-areal lateral connectivity contributes to contextual modulation, it is insufficient to replicate the high-magnitude suppression observed *in vivo*.

The second configuration incorporated a single, global V2 module to test the hypothesis of non-specific top-down feedback. While this architecture successfully attenuated L5 activity, consistent with the anatomical targeting of feedback projections to infragranular layers, it failed to suppress

Table 1: **Quantification of extra-classical surround suppression across architectures.** Normalized activity in L2/3 of the central V1 column during the eCRF protocol (center 20 Hz, surround 30 Hz), expressed as a percentage of the peak CRF response. Values closer to 0% indicate stronger suppression. The competitive V2 model (Config. 3) is the only architecture that quantitatively reproduces the $\sim 75\%$ suppression (i.e., $\sim 25\%$ residual activity) observed *in vivo* [7].

Architecture	Normalized Activity (% CRF)	Suppression Index (%)
Config. 1 (V1 lateral only)	63.32 ± 4.40	~ 36.7
Config. 2 (Linear V2 feedback)	84.07 ± 8.94	~ 16.0
Config. 3 (Competitive V2)	22.91 ± 8.32	~ 77.1
<i>In vivo</i> data [7]	~ 25	~ 75

260 L2/3. L2/3 activity remained at $84.0\% \pm 8.9\%$ of its classical receptive field baseline. Analysis of
 261 the V2 module revealed that it acted as a linear integrator, indiscriminately pooling feedforward
 262 inputs from both the central and peripheral V1 columns without undergoing internal competition.
 263 Consequently, the feedback signal functioned as an excitatory gain amplifier rather than a selective
 264 suppressive filter.

265 The third architectural configuration was the only one that achieved deep suppression in both superfi-
 266 cial and deep layers. Figure 3 presents a comprehensive comparison of the three configurations. The
 267 peri-stimulus time histograms reveal that Configurations 1 and 2 show either weak or layer-incomplete
 268 modulation. The suppression index quantifies this difference (Table 1), reaching approximately 77%
 269 only in the competitive V2 architecture.




Figure 3: **Surround Suppression is Driven by Hierarchical Competition.** (A) Diagram of the V1-V2 competitive architecture illustrating the lateral inhibitory (IL) connections. (B) Mean firing rates across V1, V2, and thalamic populations during CRF and eCRF stimulation. Only the architecture incorporating competitive lateral inhibition in V2 replicates the profound ($\sim 75\%$) suppression observed *in vivo* within V1 Layer 2/3. Statistical significance is denoted by *** ($p < 0.001$).

270 **4.3 Robust Suppression via Hierarchical Competitive Dynamics.** The integration of a hierarchi-
 271 cally segregated, competitive V2 architecture (Configuration 3) fundamentally altered the network
 272 dynamics. In this design, two distinct V2 modules, each aligned to specific feature columns in V1,
 273 interacted via strong, reciprocal inhibitory connections. Under the extra-classical receptive field
 274 stimulation protocol, activation of the peripheral column triggered a robust winner-take-all dynamic in
 275 V2. The peripheral V2 module surged in activity, actively suppressing the central V2 module, whose
 276 firing rate fell to near baseline levels, approximately 1 Hz. This selective refinement of the feedback
 277 signal profoundly modulated V1. In the central V1 microcolumn, L2/3 activity was suppressed to
 278 $22.9\% \pm 8.3\%$ of its maximum classical receptive field response, corresponding to a suppression
 279 index of approximately 77%, which quantitatively replicates the $\sim 75\%$ suppression magnitude
 280 reported in macaque V1 [7].

281 The competitive V2 architecture not only suppresses V1 activity but also profoundly silences the
 282 central V2 module itself, as shown in the detailed firing rate profiles (Table 2). To isolate the source
 283 of this modulation, we performed virtual ablation experiments. Removing the horizontal lateral
 284 connections within V1 did not abolish the surround suppression effect. This demonstrates that the

Table 2: **Detailed activity profile for the competitive V2 architecture (Config. 3).** Mean firing rates (in spikes/s $\pm$ SD) of excitatory neurons in the central V1 column and the central V2 module under CRF and eCRF stimulation with a 30 Hz peripheral stimulus. This configuration shows deep suppression in both V1 and V2 superficial layers, driven by hierarchical competition.

Area & Layer	CRF (20 Hz)	eCRF (30 Hz)
V1 L2/3	6.46 $\pm$ 0.62	2.10 $\pm$ 0.43
V1 L5	30.62 $\pm$ 0.70	28.72 $\pm$ 0.60
V2 L2/3	10.09 $\pm$ 0.95	1.11 $\pm$ 0.59
V2 L5	13.70 $\pm$ 1.83	19.98 $\pm$ 1.91

top-down feedforward-feedback loop, equipped with competitive inter-areal dynamics in higher visual cortices, is sufficient to mediate robust far-surround suppression, overcoming the spatial and temporal limitations of intra-areal lateral axons.

5 Discussion

Emergence of eCRF effects through Hierarchical Competition. Our results reveal that profound contextual modulation is an emergent property of hierarchical competitive dynamics. It is only when top-down feedback signals are non-linearly refined via lateral, winner-take-all competition within higher-order visual areas (such as V2) that the model reproduces physiological suppression depths. These findings fundamentally challenge the traditional view of visual processing as a strictly local feature-extraction pipeline. By embedding a biologically plausible laminar microcircuit within a hierarchical competitive framework, our model bridges the theoretical gap between anatomical constraints—such as the specific targeting of L2/3 and L5 by feedback projections—and the complex, global dynamics of sensory integration.

Computational Advantages of Non-linear Feedback. From a computational perspective, embedding lateral competition within the higher-order feedback loop provides an efficient mechanism for sensory filtering and predictive coding. Simple additive or multiplicative top-down modulations can lead to runaway excitation or generalized gain changes that lack spatial and feature specificity. In contrast, our simulations indicate that inter-areal competition acts as a selective spatial filter. When a visual stimulus extends beyond the classical receptive field, the competitive interactions in V2 actively sculpt the resulting feedback to V1, suppressing neural activity that encodes predictable spatial continuities.

This architecture serves two computational functions. First, it prevents the saturation of primary sensory areas during high-contrast, large-scale visual stimulation, maintaining the network in a stable asynchronous-irregular regime [29]. Second, it enhances information transmission by silencing neural populations that encode redundant contextual information, allowing the network to prioritize salient features such as boundaries, discontinuities, or unexpected stimuli [30]. This top-down competitive mechanism can therefore be understood as a circuit-level operation for sparse coding and sensory optimization, consistent with predictive coding frameworks [25, 26].

Implications for Machine Learning. The hierarchical competitive mechanism we identify offers a concrete architectural principle for the design of artificial neural networks. The core insight, that lateral inhibition in higher processing stages enables robust, context-dependent normalization, aligns with several recent trends in deep learning yet remains underexplored as an explicit design choice.

First, our winner-take-all gating of feedback bears functional similarities to attention mechanisms [27], where queries selectively weight the relevance of different spatial locations or feature channels. However, whereas standard self-attention computes pairwise similarities across all positions, incurring quadratic complexity, our lateral inhibition operates locally between modules, offering a potentially more scalable alternative for spatial normalization. This principle aligns with recent work on squeeze-and-excitation networks [15], which recalibrate channel-wise features via a lightweight gating mechanism, and could be extended to incorporate lateral competition between feature modules.

Second, the normalization dynamics we observe, in which V1 responses are scaled relative to the activity of competing V2 modules, implement a form of contrastive computation [31, 32]. Our

biologically constrained architecture suggests that hierarchical lateral inhibition could serve as an energy-efficient alternative to explicit contrastive losses, potentially enabling on-chip learning in neuromorphic hardware [33].

Third, the dynamic range compression achieved by competitive feedback offers a potential route to adversarial robustness. By normalizing responses relative to spatial context rather than absolute input magnitudes, the network becomes inherently less sensitive to perturbations that uniformly scale input intensities, a property consistent with known features of divisive normalization in biological vision [15]. Finally, our work demonstrates that these computations can be realized in fully spiking architectures, which are compatible with emerging neuromorphic platforms such as Intel’s Loihi [33] and SpiNNaker [34]. Translating the hierarchical competitive motif to trainable spiking networks, for instance via surrogate gradient methods [18], represents a promising direction for developing low-power artificial perception systems that inherit the robustness and efficiency of biological sensory processing.

Limitations and Future Directions. While our model successfully captures the core non-linear dynamics of eCRF modulation, several avenues remain for future theoretical and empirical investigation. The current implementation relies on computationally efficient, simplified point-neuron models and utilizes generalized inhibitory population pools. However, cortical inhibition is highly diverse. Future iterations will significantly benefit from disentangling specific interneuron subtypes. Furthermore, the present model operates with static synaptic weights. Incorporating spike-timing-dependent plasticity (STDP) or other biologically plausible learning rules could elucidate how these hierarchical competitive networks develop their precise topographic connectivity during visual maturation. On the practical side, currently the computational implementations of the spiking neural networks relay on CPU. Adaptations to use more efficient compute, as GPUs or neuromorphic, should improve the efficiency of this simulation, opening the door for scaling network size.

In summary, our work demonstrates that lateral inhibitory competition in higher cortical areas explains a robust surround suppression in visual areas. By grounding this principle in a large-scale spiking microcircuit, we provide a mechanistic bridge between physiological data and computational principles, offering a tangible blueprint for the next generation of robust, energy-efficient artificial perception systems.

References

- [1] John Allman, Francis Miezin, and EveLynn McGuinness. Stimulus specific responses from beyond the classical receptive field: Neurophysiological mechanisms for local-global comparisons in visual neurons. *Annual Review of Neuroscience*, 8(1):407–430, March 1985. ISSN 1545-4126. doi: 10.1146/annurev.ne.08.030185.002203. URL <http://dx.doi.org/10.1146/annurev.ne.08.030185.002203>.
- [2] Alessandra Angelucci and Jean Bullier. Reaching beyond the classical receptive field of v1 neurons: horizontal or feedback axons? *Journal of Physiology-Paris*, 97(2-3):141–154, March 2003. doi: 10.1016/j.jphysparis.2003.09.001. URL <https://doi.org/10.1016/j.jphysparis.2003.09.001>.
- [3] Wyeth Bair, James R. Cavanaugh, and J. Anthony Movshon. Time course and time-distance relationships for surround suppression in macaque v1 neurons. *The Journal of Neuroscience*, 23(20):7690–7701, August 2003. ISSN 1529-2401. doi: 10.1523/jneurosci.23-20-07690.2003. URL <http://dx.doi.org/10.1523/JNEUROSCI.23-20-07690.2003>.
- [4] Lars Schwabe, Klaus Obermayer, Alessandra Angelucci, and Paul C. Bressloff. The role of feedback in shaping the extra-classical receptive field of cortical neurons: A recurrent network model. *The Journal of Neuroscience*, 26(36):9117–9129, September 2006. ISSN 1529-2401. doi: 10.1523/jneurosci.1253-06.2006. URL <http://dx.doi.org/10.1523/JNEUROSCI.1253-06.2006>.
- [5] Alessandra Angelucci, Maryam Bijanzadeh, Lauri Nurminen, Frederick Federer, Sam Merlin, and Paul C. Bressloff. Circuits and mechanisms for surround modulation in visual cortex. *Annual Review of Neuroscience*, 40(1):425–451, July 2017. ISSN 1545-4126. doi: 10.1146/annurev-neuro-072116-031418. URL <http://dx.doi.org/10.1146/annurev-neuro-072116-031418>.

- [6] Tobias C. Potjans and Markus Diesmann. The cell-type specific cortical microcircuit: Relating structure and activity in a full-scale spiking network model. *Cerebral Cortex*, 24(3):785–806, March 2014. doi: 10.1093/cercor/bhs358. URL <https://doi.org/10.1093/cercor/bhs358>.
- [7] Laura Busse, Alex R. Wade, and Matteo Carandini. Representation of concurrent stimuli by population activity in visual cortex. *Neuron*, 64(6):931–942, December 2009. ISSN 0896-6273. doi: 10.1016/j.neuron.2009.11.004. URL <http://dx.doi.org/10.1016/j.neuron.2009.11.004>.
- [8] Paulina Anna Dąbrowska, Nicole Voges, Michael von Papen, Junji Ito, David Dahmen, Alexa Riehle, Thomas Brochier, and Sonja Grün. On the complexity of resting state spiking activity in monkey motor cortex. *Cerebral Cortex Communications*, 2(3), January 2021. ISSN 2632-7376. doi: 10.1093/texcom/tgab033. URL <http://dx.doi.org/10.1093/texcom/tgab033>.
- [9] Marc Oliver Gewaltig and Markus Diesmann. Nest (neural simulation tool). *Scholarpedia*, 2(4):1430, 2007. ISSN 1941-6016. doi: 10.4249/scholarpedia.1430. URL <http://dx.doi.org/10.4249/scholarpedia.1430>.
- [10] DC Somers, SB Nelson, and M Sur. An emergent model of orientation selectivity in cat visual cortical simple cells. *The Journal of Neuroscience*, 15(8):5448–5465, August 1995. ISSN 1529-2401. doi: 10.1523/jneurosci.15-08-05448.1995. URL <http://dx.doi.org/10.1523/jneurosci.15-08-05448.1995>.
- [11] Nobuhiko Wagatsuma, Tobias C. Potjans, Markus Diesmann, Ko Sakai, and Tomoki Fukai. Spatial and feature-based attention in a layered cortical microcircuit model. *PLoS ONE*, 8(12):e80788, December 2013. doi: 10.1371/journal.pone.0080788. URL <https://doi.org/10.1371/journal.pone.0080788>.
- [12] Jung H. Lee, Christof Koch, and Stefan Mihalas. A computational analysis of the function of three inhibitory cell types in contextual visual processing. *Frontiers in Computational Neuroscience*, 11, April 2017. ISSN 1662-5188. doi: 10.3389/fncom.2017.00028. URL <http://dx.doi.org/10.3389/fncom.2017.00028>.
- [13] Joseph Del Rosario, Stefano Coletta, Soon Ho Kim, Zach Mobbille, Kayla Peelman, Brice Williams, Alan J. Otsuki, Alejandra Del Castillo Valerio, Kendell Worden, Lou T. Blain, Lyndah Lovell, Hannah Choi, and Bilal Haider. Lateral inhibition in v1 controls neural and perceptual contrast sensitivity. *Nature Neuroscience*, 28(4):836–847, March 2025. ISSN 1546-1726. doi: 10.1038/s41593-025-01888-4. URL <http://dx.doi.org/10.1038/s41593-025-01888-4>.
- [14] Dan D Stettler, Aniruddha Das, Jean Bennett, and Charles D Gilbert. Lateral connectivity and contextual interactions in macaque primary visual cortex. *Neuron*, 36(4):739–750, November 2002. doi: 10.1016/s0896-6273(02)01029-2. URL [https://doi.org/10.1016/s0896-6273\(02\)01029-2](https://doi.org/10.1016/s0896-6273(02)01029-2).
- [15] Matteo Carandini and David J. Heeger. Normalization as a canonical neural computation. *Nature Reviews Neuroscience*, 13(1):51–62, November 2011. ISSN 1471-0048. doi: 10.1038/nrn3136. URL <http://dx.doi.org/10.1038/nrn3136>.
- [16] Maximilian Schmidt, Rembrandt Bakker, Kelly Shen, Gleb Bezgin, Markus Diesmann, and Sacha Jennifer van Albada. A multi-scale layer-resolved spiking network model of resting-state dynamics in macaque visual cortical areas. *PLOS Computational Biology*, 14(10):e1006359, October 2018. ISSN 1553-7358. doi: 10.1371/journal.pcbi.1006359. URL <http://dx.doi.org/10.1371/journal.pcbi.1006359>.
- [17] Nicolas Brunel and Xiao-Jing Wang. Effects of neuromodulation in a cortical network model of object working memory dominated by recurrent inhibition. *Journal of Computational Neuroscience*, 11(1):63–85, July 2001. ISSN 1573-6873. doi: 10.1023/a:1011204814320. URL <http://dx.doi.org/10.1023/a:1011204814320>.

- [18] Emre O. Neftci, Hesham Mostafa, and Friedemann Zenke. Surrogate gradient learning in spiking neural networks: Bringing the power of gradient-based optimization to spiking neural networks. *IEEE Signal Processing Magazine*, 36(6):51–63, November 2019. ISSN 1558-0792. doi: 10.1109/msp.2019.2931595. URL <http://dx.doi.org/10.1109/MSP.2019.2931595>.
- [19] David J. Heeger. Normalization of cell responses in cat striate cortex. *Visual Neuroscience*, 9(2):181–197, August 1992. ISSN 1469-8714. doi: 10.1017/s0952523800009640. URL <http://dx.doi.org/10.1017/s0952523800009640>.
- [20] John H. Reynolds and David J. Heeger. The normalization model of attention. *Neuron*, 61(2):168–185, January 2009. ISSN 0896-6273. doi: 10.1016/j.neuron.2009.01.002. URL <http://dx.doi.org/10.1016/j.neuron.2009.01.002>.
- [21] Kenway Louie, Mel W. Khaw, and Paul W. Glimcher. Normalization is a general neural mechanism for context-dependent decision making. *Proceedings of the National Academy of Sciences*, 110(15):6139–6144, March 2013. ISSN 1091-6490. doi: 10.1073/pnas.1217854110. URL <http://dx.doi.org/10.1073/pnas.1217854110>.
- [22] Sergey Ioffe and Christian Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift. In Francis Bach and David Blei, editors, *Proceedings of the 32nd International Conference on Machine Learning*, volume 37 of *Proceedings of Machine Learning Research*, pages 448–456, Lille, France, 07–09 Jul 2015. PMLR. URL <https://proceedings.mlr.press/v37/ioffe15.html>.
- [23] Jimmy Lei Ba, Jamie Ryan Kiros, and Geoffrey E. Hinton. Layer normalization, 2016. URL <https://arxiv.org/abs/1607.06450>.
- [24] Mengye Ren, Renjie Liao, Raquel Urtasun, Fabian H. Sinz, and Richard S. Zemel. Normalizing the normalizers: Comparing and extending network normalization schemes, 2016. URL <https://arxiv.org/abs/1611.04520>.
- [25] Rajesh P. N. Rao and Dana H. Ballard. Predictive coding in the visual cortex: a functional interpretation of some extra-classical receptive-field effects. *Nature Neuroscience*, 2(1):79–87, January 1999. doi: 10.1038/4580. URL <https://doi.org/10.1038/4580>.
- [26] Karl Friston. A theory of cortical responses. *Philosophical Transactions of the Royal Society B: Biological Sciences*, 360(1456):815–836, April 2005. ISSN 1471-2970. doi: 10.1098/rstb.2005.1622. URL <http://dx.doi.org/10.1098/rstb.2005.1622>.
- [27] Robert Desimone and John Duncan. Neural mechanisms of selective visual attention. *Annual Review of Neuroscience*, 18(1):193–222, March 1995. ISSN 1545-4126. doi: 10.1146/annurev.ne.18.030195.001205. URL <http://dx.doi.org/10.1146/annurev.ne.18.030195.001205>.
- [28] John H. Reynolds, Leonardo Chelazzi, and Robert Desimone. Competitive mechanisms subserve attention in macaque areas v2 and v4. *The Journal of Neuroscience*, 19(5):1736–1753, March 1999. ISSN 1529-2401. doi: 10.1523/jneurosci.19-05-01736.1999. URL <http://dx.doi.org/10.1523/JNEUROSCI.19-05-01736.1999>.
- [29] Nicolas Brunel. Dynamics of sparsely connected networks of excitatory and inhibitory spiking neurons. *Journal of Computational Neuroscience*, 8(3):183–208, May 2000. ISSN 1573-6873. doi: 10.1023/a:1008925309027. URL <http://dx.doi.org/10.1023/A:1008925309027>.
- [30] William E. Vinje and Jack L. Gallant. Sparse coding and decorrelation in primary visual cortex during natural vision. *Science*, 287(5456):1273–1276, February 2000. ISSN 1095-9203. doi: 10.1126/science.287.5456.1273. URL <http://dx.doi.org/10.1126/science.287.5456.1273>.
- [31] Aaron van den Oord, Yazhe Li, and Oriol Vinyals. Representation learning with contrastive predictive coding, 2018. URL <https://arxiv.org/abs/1807.03748>.
- [32] Ting Chen, Simon Kornblith, Mohammad Norouzi, and Geoffrey Hinton. A simple framework for contrastive learning of visual representations. In *International conference on machine learning*, pages 1597–1607. PmLR, 2020.

- [33] Mike Davies, Narayan Srinivasa, Tsung-Han Lin, Gautham Chinya, Yongqiang Cao, Sri Harsha Choday, Georgios Dimou, Prasad Joshi, Nabil Imam, Shweta Jain, Yuyun Liao, Chit-Kwan Lin, Andrew Lines, Ruokun Liu, Deepak Mathaikutty, Steven McCoy, Arnab Paul, Jonathan Tse, Guruguhanathan Venkataramanan, Yi-Hsin Weng, Andreas Wild, Yoonseok Yang, and Hong Wang. Loihi: A neuromorphic manycore processor with on-chip learning. *IEEE Micro*, 38(1):82–99, January 2018. ISSN 1937-4143. doi: 10.1109/mm.2018.112130359. URL <http://dx.doi.org/10.1109/MM.2018.112130359>.
- [34] Steve B. Furber, Francesco Galluppi, Steve Temple, and Luis A. Plana. The spinnaker project. *Proceedings of the IEEE*, 102(5):652–665, May 2014. ISSN 1558-2256. doi: 10.1109/jproc.2014.2304638. URL <http://dx.doi.org/10.1109/JPROC.2014.2304638>.

A Technical appendices and supplementary material

Technical appendices with additional results, figures, graphs, and proofs may be submitted with the paper submission before the full submission deadline (see above). You can upload a ZIP file for videos or code, but do not upload a separate PDF file for the appendix. There is no page limit for the technical appendices.

Note: Think of the appendix as “optional reading” for reviewers. The paper must be able to stand alone without the appendix; for example, adding critical experiments that support the main claims to an appendix is inappropriate.

A.1 Microcircuit Connectivity Matrix

Table 3: **Connection probability matrix for the canonical microcircuit.** The table shows the probability of connection between any two neurons from a source population (columns) to a target population (rows). Values are derived from Potjans and Diesmann [6]. E: Excitatory neurons; I: Inhibitory neurons.

Target \ Source	L2/3e	L2/3i	L4e	L4i	L5e	L5i	L6e	L6i
L2/3e	0.101	0.169	0.044	0.082	0.032	0.0	0.008	0.0
L2/3i	0.135	0.137	0.032	0.052	0.076	0.0	0.004	0.0
L4e	0.008	0.006	0.050	0.135	0.007	0.0	0.045	0.0
L4i	0.069	0.003	0.079	0.160	0.003	0.0	0.106	0.0
L5e	0.100	0.062	0.050	0.006	0.083	0.373	0.020	0.0
L5i	0.055	0.027	0.026	0.002	0.060	0.316	0.009	0.0
L6e	0.016	0.007	0.021	0.017	0.057	0.020	0.040	0.225
L6i	0.036	0.001	0.003	0.001	0.028	0.008	0.066	0.144

NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes], [No], or [N/A].
- [N/A] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for [N/A]).

The checklist answers are an integral part of your paper submission. They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While [Yes] is generally preferable to [No], it is perfectly acceptable to answer [No] provided a proper justification is given (e.g., error bars are not reported because it would be too computationally expensive” or “we were unable to find the license for the dataset we used”). In general, answering [No] or [N/A] is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- **Delete this instruction block, but keep the section heading “NeurIPS Paper Checklist”.**
- **Keep the checklist subsection headings, questions/answers and guidelines below.**
- **Do not modify the questions and only use the provided macros for your answers.**

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: On the Abstract we state the motivation, specify the problem, the hypothesis, explain the experimental procedure and summarized the main result of our experiments. We follow a similar strategy in the Introduction, but with a superficial explanation of the experiment and results.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: On the Discussion section we have a subsection title "Limitations and Future Directions." where we directly address the main limitations of our work.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The work does not include theoretical results.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: On the Methods section we provide the details of the simulations. And the code will be available for the reader.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No]

Justification: We plan to release the code on the next stage of review process.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.

- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Our experimental results are based on simulations of spiking neural networks, we provide all the details of the experiments and stimuli used to run the simulations.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: We report the mean and standard deviation across 10 simulation trials initialized with different random seeds, and indicate statistical significance ($p < 0.001$) in the corresponding figures.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [No]

Justification: Due to strict space constraints, the exact hardware specifications (e.g., number of CPU nodes and wall-clock time) were omitted from the main text, although the simulation framework (NEST simulator) is fully detailed to aid reproducibility.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: Our work does not include human data, any potentially harmful development, and it does not use dataset from any source.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [N/A]

Justification: Currently, our work is at the stage of simulations, so it does not imply any direct societal impact.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The work does not poses any of such risks.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: We credited all the authors, specially those of the platform used to perform the simulations.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [N/A]

Justification: This work does not introduce any asset.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.

- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: No human data was used or collected for this research.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: This work does not involve humans or animals subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: We only used LLMs for editing and formatting purposes.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.